

Handwritten Alphabet Recognition

19.04.2023

Likhith Raj DV

cse200001039@iiti.ac.in

Abhishek Jaisawal

cse200001001@iiti.ac.in

Overview

- Libraries
-
- Data Preprocessing
-
- Model1
-
- Model2
-
- Model3
-
- Image Prediction

Libraries

numpy and pandas

Pandas is used for data preprocessing and data manipulation.

Numpy is used for preprocessing a given test image for prediction.

seaborn and matplotlib

Both are used to draw plots and represent statistical data.

imblearn

Used to handle under sampling as in the dataset the number of samples in some classes outweighs the other, thus improving efficiency and performance.

tensorflow and keras

Used for building machine learning models. Keras is also used for building neural network based models.

▼ Importing libraries

```
[ ] from google.colab import drive
    drive.mount('/content/drive')
```

Mounted at /content/drive

```
▶ import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# for handling imbalancing
from imblearn.under_sampling import NearMiss
from keras.utils import np_utils

from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix

import keras
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout, BatchNormalization

import warnings
warnings.filterwarnings('ignore')
```

Data Preprocessing

Viewing the dataset

▼ Data Exploring

```
[2] df = pd.read_csv('/content/drive/MyDrive/CI/A_Z Handwritten Data.csv')
```

```
[3] df.shape
```

```
(372450, 785)
```

```
df.head()
```

	0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	...	0.639	0.640	0.641	0.642	0.643	0.644	0.645	0.646	0.647	0.648
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows × 785 columns

Dataset has 372450 samples of **28 x 28** images.

```
df.tail()
```

	0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	...	0.639	0.640	0.641	0.642	0.643	0.644	0.645	0.646	0.647	0.648
372445	25	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
372446	25	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
372447	25	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
372448	25	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
372449	25	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows × 785 columns

Handling imbalanced data

▼ Data preprocessing

```
✓ [8] # getting target variable
0s y = df['0']
    del df['0']
```

Dealing with imbalanced target

```
✓ [9] x = y.replace([0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21,22,23,24,25], ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M',
0s 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z'])
    x
```

0	A
1	A
2	A
3	A
4	A
...	..
372445	Z
372446	Z
372447	Z
372448	Z
372449	Z

Name: 0, Length: 372450, dtype: object

```
✓ [14] nM = NearMiss()
32s X_data, y_data = nM.fit_resample(df, y)
```

Replacing the integers with respective alphabets.

Here, the NearMiss algorithm is used to handle undersampling.

Normalization

Normalizing the values of images between 0 and 1 by dividing with 255. This has improved the performance of the model . Sometimes normalization leads to improved convergence rate of the machine learning algorithm during the training process.

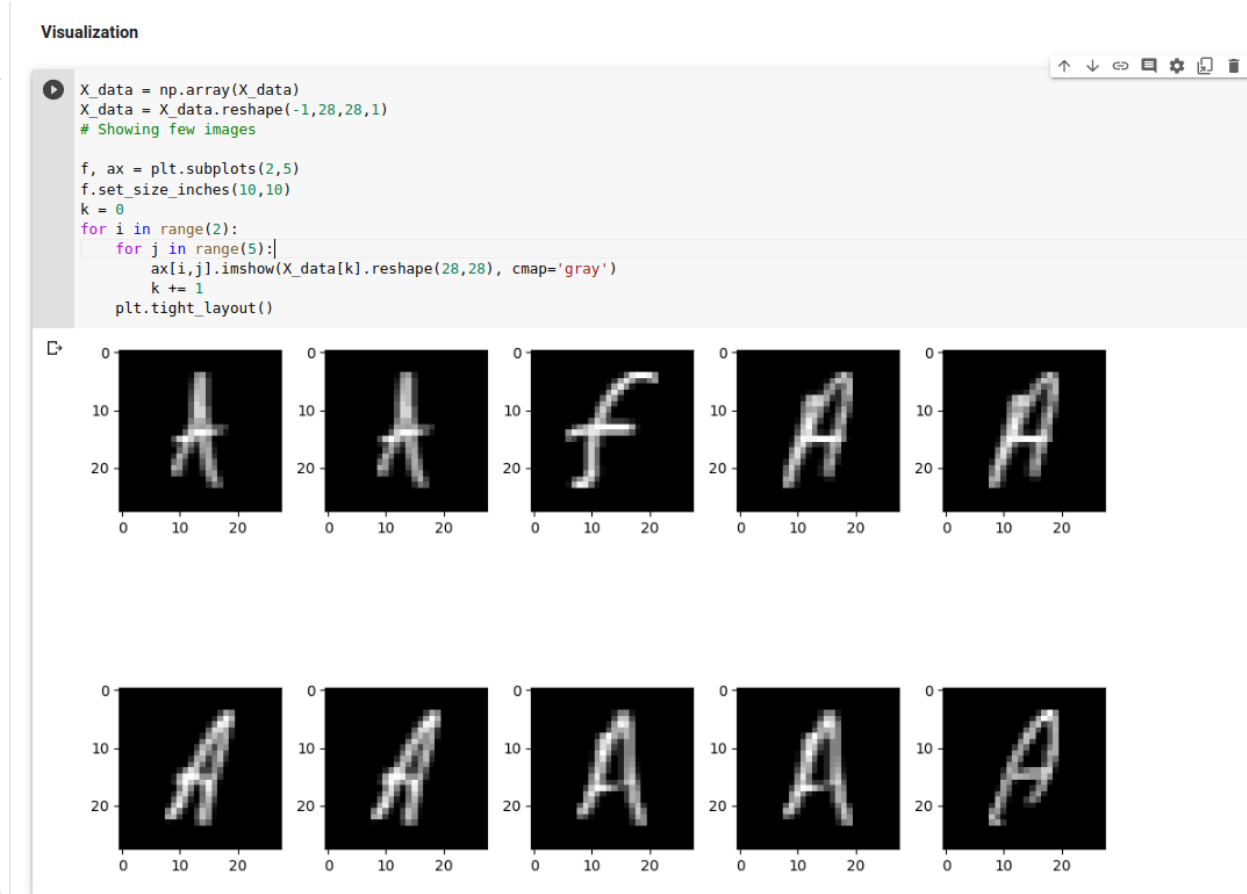
Normalization

```
[18] X_data = X_data / 255
X_data
```

	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	0.10	...	0.639	0.640	0.641	0.642	0.643	0.644	0.645	0.646	0.647	0.648
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
...
29115	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29116	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29117	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29118	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29119	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

29120 rows × 784 columns

Visualizing data



Splitting into Test and Train data.

Here, we are splitting the given dataset into train and test data by splitting it in the ratio of 80% : 20%.

Train test split

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X_data, y, test_size=0.2, random_state=102)
```

```
[ ] X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
((23296, 28, 28, 1), (5824, 28, 28, 1), (23296, 26), (5824, 26))
```


Model1

Architecture and Summary

Model 1

```
#Build an ordinary "Deep Learning" model with CNN and maxpooling by using Keras.
model = Sequential()
model.add(Conv2D(32, (5, 5), input_shape=(28, 28, 1), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
#Choose an optimizer and compile the model.
model.compile(optimizer = Adam(learning_rate = 0.01), loss = 'categorical_crossentropy', metrics = ['accuracy'])
#And print the summary of the model.
print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 24, 24, 32)	832
max_pooling2d (MaxPooling2D)	(None, 12, 12, 32)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 128)	589952
dense_1 (Dense)	(None, 26)	3354
Total params: 594,138		
Trainable params: 594,138		
Non-trainable params: 0		
None		

This is the first model we have implemented. It has five layers and a total of 594138 trainable parameters.

Training Model1

```
[ ] history = model.fit(X_train,y_train,epochs=15, batch_size=128, validation_data=(X_test,y_test))

Epoch 1/15
182/182 [=====] - 13s 6ms/step - loss: 0.4713 - accuracy: 0.8688 - val_loss: 0.1716 - val_accuracy: 0.9511
Epoch 2/15
182/182 [=====] - 1s 4ms/step - loss: 0.1152 - accuracy: 0.9664 - val_loss: 0.1535 - val_accuracy: 0.9511
Epoch 3/15
182/182 [=====] - 1s 4ms/step - loss: 0.0731 - accuracy: 0.9778 - val_loss: 0.1231 - val_accuracy: 0.9511
Epoch 4/15
182/182 [=====] - 1s 4ms/step - loss: 0.0495 - accuracy: 0.9847 - val_loss: 0.1096 - val_accuracy: 0.9511
Epoch 5/15
182/182 [=====] - 1s 4ms/step - loss: 0.0400 - accuracy: 0.9879 - val_loss: 0.1285 - val_accuracy: 0.9511
Epoch 6/15
182/182 [=====] - 1s 4ms/step - loss: 0.0403 - accuracy: 0.9883 - val_loss: 0.1109 - val_accuracy: 0.9511
Epoch 7/15
182/182 [=====] - 1s 4ms/step - loss: 0.0491 - accuracy: 0.9849 - val_loss: 0.1449 - val_accuracy: 0.9511
Epoch 8/15
182/182 [=====] - 1s 4ms/step - loss: 0.0464 - accuracy: 0.9867 - val_loss: 0.1511 - val_accuracy: 0.9511
Epoch 9/15
182/182 [=====] - 1s 4ms/step - loss: 0.0340 - accuracy: 0.9910 - val_loss: 0.1843 - val_accuracy: 0.9511
Epoch 10/15
182/182 [=====] - 1s 5ms/step - loss: 0.0405 - accuracy: 0.9880 - val_loss: 0.1539 - val_accuracy: 0.9511
Epoch 11/15
182/182 [=====] - 1s 6ms/step - loss: 0.0297 - accuracy: 0.9911 - val_loss: 0.1859 - val_accuracy: 0.9511
Epoch 12/15
182/182 [=====] - 1s 6ms/step - loss: 0.0362 - accuracy: 0.9903 - val_loss: 0.1874 - val_accuracy: 0.9511
Epoch 13/15
182/182 [=====] - 1s 5ms/step - loss: 0.0314 - accuracy: 0.9916 - val_loss: 0.1606 - val_accuracy: 0.9511
Epoch 14/15
182/182 [=====] - 1s 4ms/step - loss: 0.0285 - accuracy: 0.9927 - val_loss: 0.2444 - val_accuracy: 0.9511
Epoch 15/15
182/182 [=====] - 1s 4ms/step - loss: 0.0388 - accuracy: 0.9907 - val_loss: 0.2313 - val_accuracy: 0.9511
```

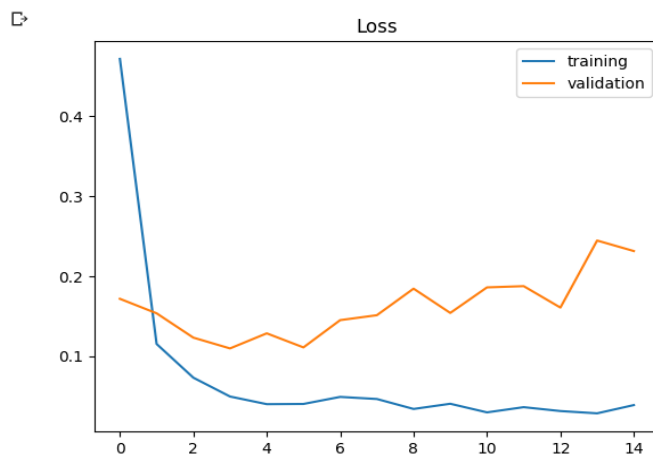
Accuracy = 99%

Loss = 0.0388

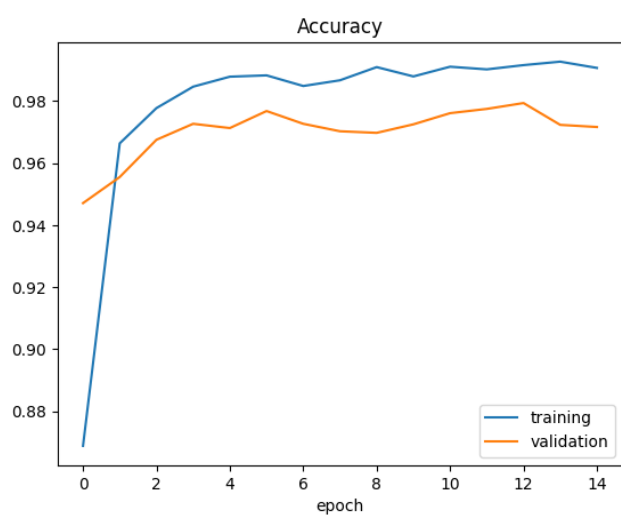
Plots for loss vs epoch

```
[ ] model.save('/content/drive/MyDrive/CI/model1.h5')
```

```
plt.figure(1)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend(['training', 'validation'])
plt.title('Loss')
plt.xlabel('epoch')
plt.figure(2)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'])
plt.title('Accuracy')
plt.xlabel('epoch')
plt.show()
```



Plots for accuracy vs epoch



Model2

Architecture and Summary

Model 2

```

model2 = Sequential()

model2.add(Conv2D(64, (5, 5), input_shape=(28, 28, 1), activation='relu', padding='same'))
model2.add(Conv2D(32, (5, 5), input_shape=(28, 28, 1), activation='relu', padding='same'))
model2.add(MaxPooling2D(pool_size=(2, 2)))

model2.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model2.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model2.add(MaxPooling2D(pool_size=(2, 2)))

model2.add(Dropout(0.2))

model2.add(Flatten())
model2.add(Dense(128, activation='relu'))
model2.add(Dense(num_classes, activation='softmax'))

model2.compile(loss='mean_squared_error', optimizer='adam', metrics=['accuracy'])
print(model2.summary())

```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d_33 (Conv2D)	(None, 28, 28, 64)	1664
conv2d_34 (Conv2D)	(None, 28, 28, 32)	51232
max_pooling2d_17 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_35 (Conv2D)	(None, 14, 14, 128)	36992
conv2d_36 (Conv2D)	(None, 14, 14, 128)	147584
max_pooling2d_18 (MaxPooling2D)	(None, 7, 7, 128)	0
dropout_8 (Dropout)	(None, 7, 7, 128)	0
flatten_9 (Flatten)	(None, 6272)	0
dense_18 (Dense)	(None, 128)	802944
dense_19 (Dense)	(None, 26)	3354

Total params: 1,043,770
 Trainable params: 1,043,770
 Non-trainable params: 0

This is the first model we have implemented. It has five layers and a total of 1,043,770 trainable parameters.

Training Model2

```
history = model2.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, batch_size=256, verbose=2)
```

Epoch	91/91	loss	accuracy	val_loss	val_accuracy	time/epoch	time/step
Epoch 1/10	91/91	0.0201	0.5848	0.0067	0.8838	6s/epoch	70ms/step
Epoch 2/10	91/91	0.0056	0.9001	0.0043	0.9220	3s/epoch	34ms/step
Epoch 3/10	91/91	0.0041	0.9247	0.0035	0.9341	3s/epoch	32ms/step
Epoch 4/10	91/91	0.0033	0.9381	0.0032	0.9366	3s/epoch	31ms/step
Epoch 5/10	91/91	0.0029	0.9431	0.0030	0.9394	3s/epoch	33ms/step
Epoch 6/10	91/91	0.0026	0.9472	0.0029	0.9411	3s/epoch	31ms/step
Epoch 7/10	91/91	0.0024	0.9507	0.0027	0.9433	3s/epoch	33ms/step
Epoch 8/10	91/91	0.0023	0.9513	0.0028	0.9425	3s/epoch	30ms/step
Epoch 9/10	91/91	0.0023	0.9513	0.0027	0.9447	3s/epoch	30ms/step
Epoch 10/10	91/91	0.0021	0.9539	0.0028	0.9404	3s/epoch	32ms/step

Accuracy = 95%

Loss = 0.95

Plots for loss vs epoch and accuracy vs epoch

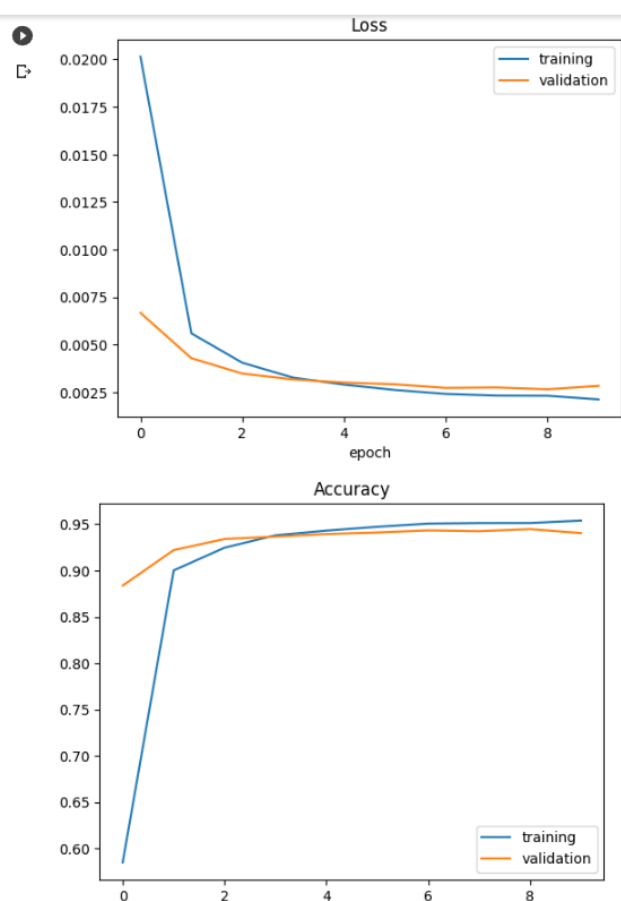


Image Prediction

```

testimage.py
import pygame
import tensorflow as tf
from tensorflow import keras
import numpy as np
import cv2

def predict_alphabet():
    # Load the Keras model
    model = keras.models.load_model('model.h5')

    # Load the image
    image = cv2.imread('alphabet.png', 0)
    image = cv2.resize(image, (28, 28)) # Resize the image to the input shape of the model

    # Preprocess the image
    image = np.expand_dims(image, axis=0) # Add batch dimension
    image = image / 255.0 # Normalize the pixel values to be between 0 and 1

    # Make predictions on the image
    predictions = model.predict(image)

    # Print the predicted class
    predicted_class = np.argmax(predictions)
    return predicted_class

def start_game():
    pygame.init()

    # Set up the font
    font = pygame.font.SysFont('Arial', 30)

    # Render the text

    # Set the size of the window
    screen = pygame.display.set_mode((600, 600))
    screen.fill((0, 0, 0))

    # Set the title of the window
    pygame.display.set_caption("Press S to Guess and C to Clear")

    # Set the default color to white
    color = (255, 255, 255)

    # Set the thickness of the brush
    brush_size = 10.75

    # Set a boolean variable to indicate if the mouse button is pressed
    drawing = False

    # Start the main loop
    while True:
        # Check for events
        for event in pygame.event.get():
            # Quit the program if the window is closed
            if event.type == pygame.QUIT:
                pygame.quit()
                quit()

            # Set drawing to True if the mouse button is pressed
            elif event.type == pygame.MOUSEBUTTONDOWN:
                drawing = True

            # Set drawing to False if the mouse button is released
            elif event.type == pygame.MOUSEBUTTONUP:
                drawing = False

            # Draw a line if the mouse button is pressed and the mouse is moving
            elif event.type == pygame.MOUSEMOTION and drawing:
                pygame.draw.circle(screen, color, event.pos, brush_size)

        # Update the display
        pygame.display.update()

        # Save the image if the S key is pressed
        keys = pygame.key.get_pressed()
        if keys[pygame.K.s]:
            surface = pygame.display.get_surface()
            image = pygame.transform.scale(surface, (32, 32))
            global ind
            pygame.image.save(image, "alphabet.png")
            print("Image saved")
            ind = predict_alphabet()
            pygame.time.delay(500)
            mystr = "The alphabet is " + chr(ord('A')+ind)
            text = font.render(mystr, True, (255, 255, 255))
            text_rect = text.get_rect(center=(200, 200))
            screen.blit(text, text_rect)
            print(mystr)
            pygame.display.update()

        # Clear screen
        if keys[pygame.K.c]:
            screen.fill((0, 0, 0))

    start_game()

```

A pygame based app to take image input and then test the input.


```
from tensorflow import keras
import numpy as np
import cv2

def predict_alphabet
    # Load the Keras
    model = keras.mo

    # Load the image
    image = cv2.imre
    image = cv2.resi

    # Preprocess the
    image = np.expan
    image = image /

    # Make predictio
    predictions = mo

    # Print the pred
    predicted_class =
```

Press S to Guess and C to clear

The alphabet is X

MS 5 OUTPUT DEBUG C

04-19 02:01:04.352150. in tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT

saved

=====] - 0s 111ms/step

alphabet is X



Thank You!